

# Today's Topics

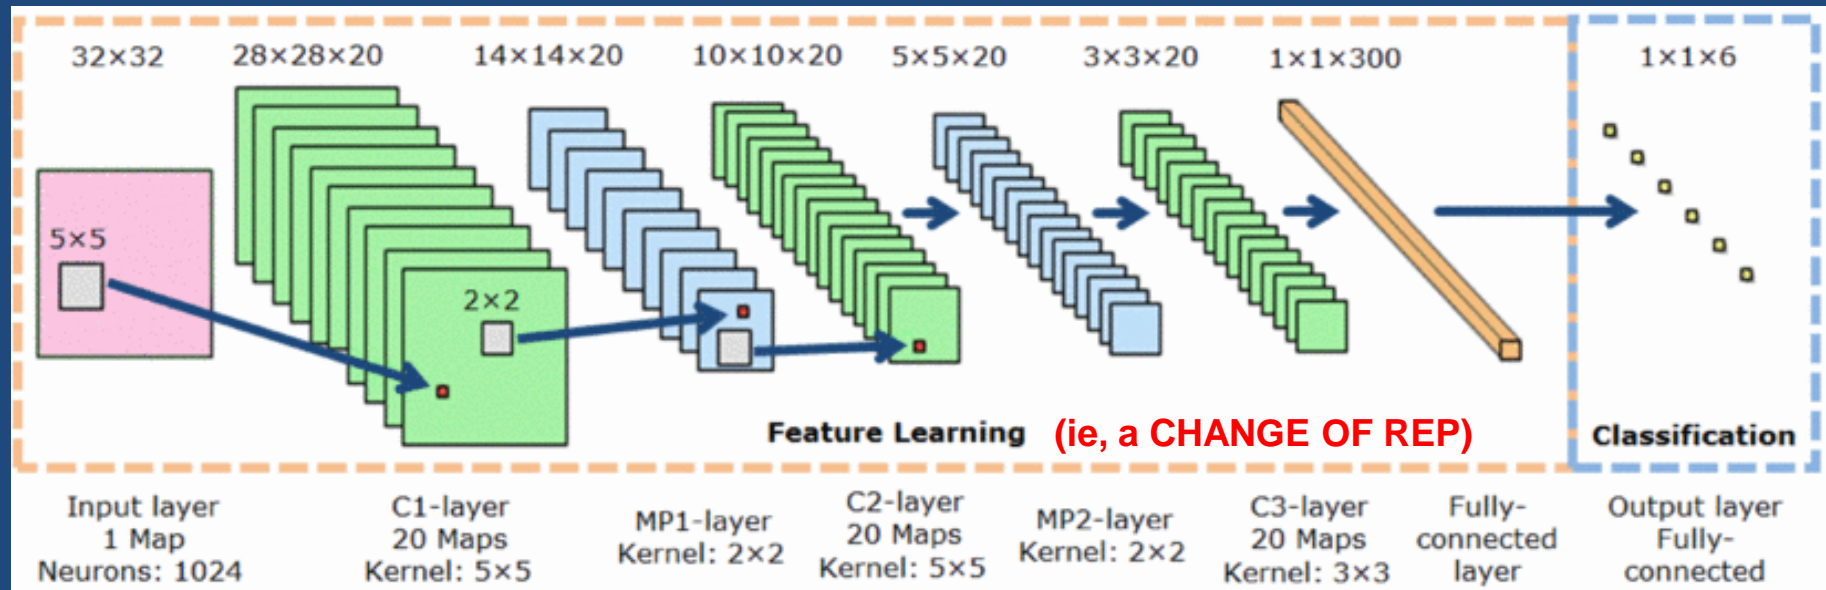


- Some Lab 3 Tips and Discussion
- Still Need a Lab 3 Partner?
- If Still Working on Lab 2, Complete ASAP!
- Some Advanced Weight-Learning Topics (by Yujia Bao)
- Interpreting ANN Outputs as Probability Distributions and the Cross-Entropy Error Function (might be covered next week)
- Next week: An Overview of Some Deep ML Projects at American Family

# Back to Deep ANNs

## - Convolution & Max Pooling (Repeat)


C = Convolution, MP = Max Pooling



My implementation (no dropout yet) of the above topology takes about 1-2 mins per epoch on the provided TRAIN/TUNE/TEST set (I measure TRAIN, TUNE and TEST accuracy after each epoch)


# Details: Image to First CONV

- There are MULTIPLE, *INDEPENDENT* plates (so each can learn a different ‘derived feature’)
- *Within* a plate, there is only one set of weights, shared for all ‘starting positions’




# Details: Image to First CONV

- There are MULTIPLE, *INDEPENDENT* plates (so each can learn a different ‘derived feature’)
- *Within* a plate, there is only one set of weights, shared for all ‘starting positions’



Image

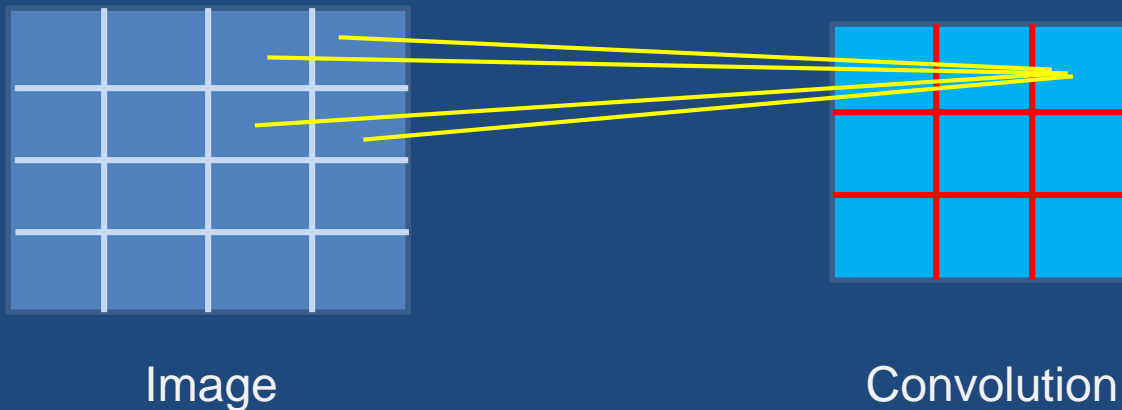


Convolution

**Note that  
the ‘sliding  
windows’  
OVERLAP  
(a design  
choice)**


# Details: Image to First CONV

- There are MULTIPLE, *INDEPENDENT* plates (so each can learn a different ‘derived feature’)
- *Within* a plate, there is only one set of weights, shared for all ‘starting positions’




# Details: Image to First CONV

- There are MULTIPLE, *INDEPENDENT* plates (so each can learn a different ‘derived feature’)
- *Within* a plate, there is only one set of weights, shared for all ‘starting positions’



# CONV Layer Weight Sharing

- The SAME WEIGHTs are used for all 'sliding window' locations per plate
- You will only backprop through CONV HUs that are the MAX in some POOL window
- You might want to scale (ie, divide)  $\eta$  by the number of cells in the following POOL layer (ie, the number of MAX'es computed)



**function** BACK-PROP-LEARNING(*examples*, *network*) **returns** a neural network

**inputs:** *examples*, a set of examples, each with input vector  $\mathbf{x}$  and output vector  $\mathbf{y}$   
*network*, a multilayer network with  $L$  layers, weights  $w_{i,j}$ , activation function  $g$

**local variables:**  $\Delta$ , a vector of errors, indexed by network node

**repeat**

**for each** weight  $w_{i,j}$  **in** *network* **do**  
     $w_{i,j} \leftarrow$  a small random number

**for each** example  $(\mathbf{x}, \mathbf{y})$  **in** *examples* **do**  
    /\* Propagate the inputs forward to compute the outputs \*/  
    **for each** node  $i$  in the input layer **do**  
       $a_i \leftarrow x_i$

**for**  $\ell = 2$  **to**  $L$  **do**  
      **for each** node  $j$  in layer  $\ell$  **do**  
         $in_j \leftarrow \sum_i w_{i,j} a_i$   
         $a_j \leftarrow g(in_j)$

    /\* Propagate deltas backward from output layer to input layer \*/  
    **for each** node  $j$  in the output layer **do**  
       $\Delta[j] \leftarrow g'(in_j) \times (y_j - a_j)$

**for**  $\ell = L - 1$  **to**  $1$  **do**  
      **for each** node  $i$  in layer  $\ell$  **do**  
         $\Delta[i] \leftarrow g'(in_i) \sum_j w_{i,j} \Delta[j]$

    /\* Update every weight in network using deltas \*/  
    **for each** weight  $w_{i,j}$  in *network* **do**  
       $w_{i,j} \leftarrow w_{i,j} + \alpha \times a_i \times \Delta[j]$

**until** some stopping criterion is satisfied

**return** *network*

**Figure 18.24** The back-propagation algorithm for learning in multilayer networks.

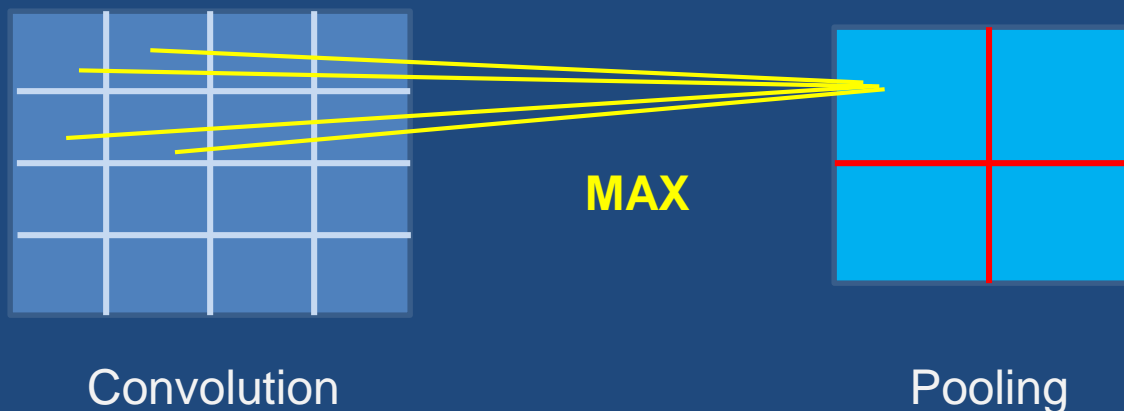
the weight-update rule for the weights between the inputs and the hidden layer is identical to the update rule for the output layer.

From CS 540  
textbook, by  
Russel and  
Norvig, 3<sup>rd</sup> ed



# Details: $N^{\text{th}}$ CONV to $N^{\text{th}}$ POOL

- I decided to have one POOL per CONV
- POOLS compute MAX via code (no wghts)
- In general, could have several POOLS of different sizes per CONV



# Details: $N^{\text{th}}$ CONV to $N^{\text{th}}$ POOL


- I *decided* to have one POOL per CONV
- POOLs compute MAX via code (no wghts)
- In general, could have several POOLs of *different sizes* per CONV



**Note that  
the 'sliding  
windows'  
DO NOT  
OVERLAP  
(a design  
choice)**


# Details: $N^{\text{th}}$ POOLS to $N^{\text{th}}$ CONV

- We want a CONV layer to look at ALL nodes at previous POOL layer, but only look at the same window
- This allows learning combinations of derived features



# Details: $N^{\text{th}}$ POOLS to $N^{\text{th}}$ CONV


- We want a CONV layer to look at ALL nodes at previous POOL layer, but only look at the same window
- This allows learning combinations of derived features




# Details: N<sup>th</sup> CONV to 1 HU Layer

- We want the flat layer of HUs to look all CONV/POOL nodes at final CONV/POOL layer.
- Initially Lab 3 said to use the following, but ok to follow Slide 2's drawing and have a 3<sup>rd</sup> CONV layer (I did this)

## CONV-POOL-CONV-POOL-FLAT-OUTPUT



# BP'ing Through POOL Nodes



- My implementation only BPs through the CONV nodes that are the MAX for some POOL window (all others have deviation=0)
- There are no learnable weights (and biases) between CONV and POOL layers (where needed I assume  $wgt=1$ , but never change it – ie, the MAX calc is not done via wgt'ed sums and an activation function; it is just Java code)
- POOL nodes compute 'deviations' (ie, the intermediate calc in the BP algorithm), but don't change wgt's and biases
- If POOL windows do not overlap (and only 1 POOL plate per CONV plate), a CONV node only sums the deviations from at most one POOL node in the next layer

# Pictorially

